


TAKEAWAY

The factory made hundreds of model comparisons affordable and auditable—then gave collaborators enough control to disprove results that initially looked convincing.


MEASURE THE SPEEDUP



Split by cost

On the Yoruba corpus, reading, fitting the BPE, encoding 260M characters and moving the flat token store to GPU took 53 seconds. One 98M-parameter training run took 85 minutes—a 96× ratio that defines the boundary between interactive work and unattended queues.

INTERACTIVE → UNATTENDED



THE FACTORY LOOP

01 PREPARE

tokenize once; fingerprint

02 ESTIMATE

measure this GPU

03 PRETRAIN

one cell or a fleet

04 READ

curves and result records


THE RESULT

The claims gate reports 6 supported, 2 unsupported and 1 underpowered claim. Failures stay visible.

THE STANDARD

A result must be reproducible enough for another person to challenge and overturn it. That happened twice—and is evidence the tooling worked.

NINE-FUNCTION SURFACE



Records survive

Every setting that can move a number belongs in the record. Corpora and vocabularies carry fingerprints; run reuse refuses mismatched settings. The current evidence base contains 197 pretraining runs, 278 fine-tuning records and 892 individual fine-tuning runs.

ONE OPERATIONAL VIEW

